

Unit 9

Neural Networks

EE-UY 4563/EL-GY 9123: INTRODUCTION TO MACHINE LEARNING
PROF. SUNDEEP RANGAN



Learning Objectives


- ❑ Mathematically describe a neural network with a single hidden layer
 - Describe mappings for the hidden and output units
- ❑ Manually compute output regions for very simple networks
- ❑ Select the loss function based on the problem type
- ❑ Build and train a simple neural network in Keras
- ❑ Write the formulas for gradients using backpropagation
- ❑ Describe mini-batches in stochastic gradient descent

Outline

 Motivating Idea: Nonlinear classifiers from linear features

- ☐ Neural Networks
- ☐ Neural Network Loss Function
- ☐ Stochastic Gradient Descent
- ☐ Building and Training a Network in Keras
 - Synthetic data
 - In-class exercise
 - MNIST
- ☐ Backpropagation Training


Most Datasets are not Linearly Separable



- Consider simple synthetic data
 - See figure to the left
 - 2D features
 - Binary class label
- Not separated linearly


All code in <https://github.com/sdrangan/introml/blob/master/neural/synthetic.ipynb>

From Linear to Nonlinear




- Idea: Build nonlinear region from linear decisions
- Possible form for a classifier:
 - Step 1: Classify into small number of linear regions
 - Step 2: Predict class label from step 1 decisions

Step 1 Outputs and Step 2 Outputs



- ❑ Each output from step 1 is from a linear classifier with soft decision
 - Similar to logistic regression
- ❑ Final output is a weighted average of step 1 outputs using the weights
 - Weights are indicated on top of the figures

A Possible Two Stage Classifier



□ Input sample: $\mathbf{x} = (x_1, x_2)^T$

□ First step: **Hidden layer**

- Take $N_H = 4$ linear discriminants

$$z_{H,1} = \mathbf{w}_{H,1}^T \mathbf{x} + b_{H,1}$$

\vdots

$$z_{H,N_H} = \mathbf{w}_{H,M}^T \mathbf{x} + b_{H,M}$$

- Make a soft decision on each linear region

$$u_{H,m} = g(z_{H,m}) = 1/(1 + e^{-z_{H,m}})$$


□ Second step: **Output layer**

- Linear step $z_O = \mathbf{w}_O^T \mathbf{u}_H + b_O$
- Soft decision: $u_O = g(z_O)$

Model Block Diagram

□ Hidden layer: $\mathbf{z}_H = \mathbf{W}_H^T \mathbf{x} + \mathbf{b}_H$, $\mathbf{u}_H = g(\mathbf{z}_H)$

□ Output layer: $\mathbf{z}_O = \mathbf{W}_O^T \mathbf{u}_H + \mathbf{b}_O$, $u_O = g(\mathbf{z}_O)$



Training the Model

□ Model in matrix form:

- Hidden layer: $\mathbf{z}_H = \mathbf{W}_H^T \mathbf{x} + \mathbf{b}_H$, $\mathbf{u}_H = g(\mathbf{z}_H)$
- Output layer: $z_O = \mathbf{W}_O^T \mathbf{u}_H + \mathbf{b}_O$, $u_O = g(z_O)$

□ $z_O = F(\mathbf{x}, \theta)$: Linear output from final stage

- Parameters: $\theta = (\mathbf{W}_H, \mathbf{W}_O, b_H, b_O)$

□ Get training data $(\mathbf{x}_i, y_i), i = 1, \dots, N$

□ Define loss function: $L(\theta) := -\sum_{i=1}^N \ln P(y_i | x_i, \theta)$,

□ Pick parameters to minimize loss:

$$\hat{\theta} = \arg \min_{\theta} L(\theta)$$


- Will discuss how to do this minimization later

Results


□ Neural network finds a nonlinear region

□ Plot shows:

- Blue circles: Negative samples
- Green circles: Positive samples
- Red color: Classifier soft probability $g(z_o)$




Hidden Layer outputs and final outputs



- ❑ Each hidden node is a linear classifier with soft decision (Logistic regression)
- ❑ Final output is a weighted average of step 1 outputs using the weights indicated on top of the figures

Outline

❑ Motivating Idea: Nonlinear classifiers from linear features

 ❑ Neural Networks

❑ Neural Network Loss Function

❑ Stochastic Gradient Descent

❑ Building and Training a Network in Keras

- Synthetic data
- MNIST

❑ Backpropagation Training

General Structure

□ **Input:** $\mathbf{x} = (x_1, \dots, x_d)$

- N_I = number of features


□ **Hidden layer:**

- Linear transform: $\mathbf{z}_H = \mathbf{W}_H^T \mathbf{x} + \mathbf{b}_H$
- Activation function: $\mathbf{u}_H = g_{act}(\mathbf{z}_H)$
- Dimension: N_H hidden units

□ **Output layer:**


- Linear transform: $\mathbf{z}_O = \mathbf{W}_O^T \mathbf{u}_H + \mathbf{b}_O$
- Output function: $\mathbf{u}_O = g_{out}(\mathbf{z}_O)$
- Dimension: N_O = number of classes / outputs

□ Can be used for classification or regression, with different decision functions




General Neural Net Block Diagram

- Hidden layer: $\mathbf{z}_H = \mathbf{W}_H^T \mathbf{x} + \mathbf{b}_H$, $\mathbf{u}_H = g_{act}(\mathbf{z}_H)$
- Output layer: $\mathbf{z}_O = \mathbf{W}_O^T \mathbf{u}_H + \mathbf{b}_O$
- Response map: $\hat{y} = u_O = g_{out}(\mathbf{z}_O)$ (for the case of binary classes)



Terminology

- ❑ **Hidden variables:** the variables $\mathbf{z}_H, \mathbf{u}_H$
 - These are not directly observed
- ❑ **Hidden units:** The functions that compute:
 - $z_{H,i} = \sum_j W_{H,ji}x_j + b_{H,i}$, $u_{H,i} = g(z_{H,i})$
 - The function $g(z)$ called the **activation function**
- ❑ **Output units:** The functions that compute
 - $z_{O,i} = \sum_j W_{O,ji}u_{H,j} + b_{O,i}$



Response Map or Output Activation


- ❑ Last layer depends on type of response
- ❑ Binary classification: $y = \pm 1$
 - z_O is a scalar
 - Hard decision: $\hat{y} = \text{sign}(z_O)$
 - Soft decision: $P(y = 1|x) = 1/(1 + e^{-z_O})$
- ❑ Multi-class classification: $y = 1, \dots, K$
 - $\mathbf{z}_O = [z_{O,1}, \dots, z_{O,K}]^T$ is a vector
 - Hard decision: $\hat{y} = \arg \max_k z_{O,k}$
 - Soft decision: $P(y = k|x) = S_k(\mathbf{z}_O)$, $S_k(\mathbf{z}_O) = \frac{e^{z_{O,k}}}{\sum_{\ell} e^{z_{O,\ell}}}$ (softmax)
- ❑ Regression: $\mathbf{y} \in R^K$
 - $\hat{\mathbf{y}} = \mathbf{z}_O$ (linear output layer)

Hidden Activation Function

Two common activation functions


Sigmoid:

- $g_{act}(z) = 1/(1 + e^{-z})$
- Benefits: Values are bounded
- Often used for small networks



Rectified linear unit (ReLU):

- $g_{act}(z) = \max(0, z)$
- Can add sparsity (more on this later)
- Often used for larger networks
- Esp. in combination with dropout



Number of Parameters

| Layer | Parameter | Symbol | Number parameters |
|--------------|-----------|--------|-------------------------------|
| Hidden layer | Bias | b_H | N_H |
| | Weights | W_H | $N_H N_I$ |
| Output layer | Bias | b_O | N_O |
| | Weights | W_O | $N_O N_H$ |
| Total | | | $N_H(N_I + 1) + N_O(N_H + 1)$ |



□ Sizes:

- N_I = input dimension, N_H = number of hidden units, N_O = output dimension

□ N_H = number of hidden units is a free parameter

□ Discuss selection later

Inspiration from Biology



Simple model of neurons

- Dendrites: Input currents from other neurons
- Soma: Cell body, accumulation of charge
- Axon: Outputs to other neurons
- Synapse: Junction between neurons

Operation:

- Take weighted sum of input current
- Outputs when sum reaches a threshold


Each neuron is like one unit in neural network

History

- ❑ Interest in understanding the brain for thousands of years
- ❑ 1940s: Donald Hebb. Hebbian learning for neural plasticity
 - Hypothesized rule for updating synaptic weights in biological neurons
- ❑ 1950s: Frank Rosenblatt: Coined the term perceptron
 - Essentially single layer classifier, similar to logistic classification
 - Early computer implementations
 - But, Limitations of linear classifiers and computer power
- ❑ 1960s: Backpropagation: Efficient way to train multi-layer networks
 - More on this later
- ❑ 1980s: Resurgence with greater computational power
- ❑ 2005+: Deep networks
 - Many more layers. Increased computational power and data
 - Enabled first breakthroughs in various image and text processing.
 - Next lecture



Outline

- ❑ Motivating Idea: Nonlinear classifiers from linear features
- ❑ Neural Networks
-  ❑ Neural Network Loss Function
- ❑ Stochastic Gradient Descent
- ❑ Building and Training a Network in Keras
 - Synthetic data
 - MNIST
- ❑ Backpropagation Training

Training a Neural Network

- Given data: $(x_i, y_i), i = 1, \dots, N$
- Learn parameters: $\theta = (W_H, b_H, W_o, b_o)$
 - Weights and biases for hidden and output layers

- Will minimize a **loss function**: $L(\theta)$

$$\hat{\theta} = \arg \min_{\theta} L(\theta)$$

- $L(\theta)$ = measures how well parameters θ fit training data (x_i, y_i)

Note on Indexing

- ❑ Neural networks are often processed in **batches**
 - Set of training or test samples
- ❑ Need notation for single and batch input case
- ❑ For a **single** input x
 - x_j = j-th feature of the input
 - $z_{H,j}, u_{H,j}, z_{O,j}$ = j-th component of hidden and output variables
 - H and O stand for Hidden and Output. Not an index
 - Write x, z_O, y if they are scalar (i.e. do not write index)
- ❑ For a **batch** of inputs x_1, \dots, x_N
 - x_{ij} = j-th feature of the input sample i
 - $z_{H,ij}, u_{H,ij}, z_{O,ij}$ = j-th component of hidden and output variables for sample i

Selecting the Right Loss Function

- Depends on the problem type
- Always compare final output z_{oi} with target y_i

| Problem | Target y_i | Output z_{oi} | Loss function | Formula |
|--------------------------------|---------------------------------|--|---------------------------|--|
| Regression | $y_i = \text{Scalar real}$ | $z_{oi} = \text{Prediction of } y_i$ Scalar output / sample | Squared / L2 loss | $\sum_i (y_i - z_{oi})^2$ |
| Regression with vector samples | $y_i = (y_{i1}, \dots, y_{iK})$ | $z_{oik} = \text{Prediction of } y_{ik}$ K outputs / sample | Squared / L2 loss | $\sum_{ik} (y_{ik} - z_{oik})^2$ |
| Binary classification | $y_i = \{0,1\}$ | $z_{oi} = \text{"logit" score}$ Scalar output / sample | Binary cross entropy | $\sum_i -y_i z_{oi} + \ln(1 + e^{y_i z_i})$ |
| Multi-class classification | $y_i = \{1, \dots, K\}$ | $z_{oik} = \text{"logit" scores}$ K outputs / sample | Categorical cross entropy | $\sum_i \ln \left(\sum_k e^{z_{oik}} \right) - \sum_k r_{ik} z_{oik}$ |

Loss Function: Regression

□ Regression case:

- y_i = scalar target variable for sample i
- Typically continuous valued

□ Output layer:

- z_{oi} = estimate of y_i

□ Loss function: Use L2 loss

$$L(\theta) = \sum_{i=1}^N (y_i - z_{oi})^2$$

□ For vector $\mathbf{y}_i = (y_{i1}, \dots, y_{iK})$, use vector L2 loss

$$L(\theta) = \sum_{i=1}^N \sum_{j=1}^K (y_{ik} - z_{oik})^2$$

Loss Function: Binary Classification

□ Binary classification: $y_i = \{0,1\}$ = class label

□ Loss function = negative log likelihood

$$L(\theta) = - \sum_{i=1}^N \ln P(y_i | x_i, \theta), \quad P(y_i = 1 | x_i, \theta) = \frac{1}{1 + e^{-z_{oi}}}$$

- Output z_{oi} called the **logit score**
- z_{oi} scalar.

□ From lecture on logistic regression:

$$\ln P(y_i | x_i, \theta) = y_i u_i + (1 - y_i) \ln(1 - u_i) = y_i z_{oi} - \ln[1 + e^{y_i z_{oi}}]$$
$$u_i = P(y_i = 1 | x_i, \theta)$$

- Called the **binary cross-entropy**

Loss Function: Multi-Class Classification 1

□ $y_i = \{1, \dots, K\}$ = class label

□ Output: $\mathbf{z}_{oi} = (z_{o,i1}, \dots, z_{o,iK})$

- K outputs. One per class
- Also called the **logit score**

□ Likelihood given by **softmax**:

$$P(y_i = k | \mathbf{x}_i, \theta) = g_k(\mathbf{z}_{oi}), \quad g_k(\mathbf{z}_{oi}) = \frac{e^{z_{o,ik}}}{\sum_{\ell} e^{z_{o,i\ell}}}$$

- Assigns class highest probability with highest logit score

Loss Function: Multi-Class Classification 2

□ $y_i = \{1, \dots, K\}$ = class label

□ Define **one-hot** coded response

$$r_{ik} = \begin{cases} 1 & y_i = k \\ 0 & y_i \neq k \end{cases}$$

◦ $\mathbf{r}_i = (r_{i1}, \dots, r_{iK})$ is K -dimensional


□ Negative log-likelihood given by:

$$\ln P(y_i = k | \mathbf{x}_i, \theta) = \sum_{k=1}^K r_{ik} \ln P(y_i = k | \mathbf{x}_i, \theta), \quad P(y_i = k | \mathbf{x}_i, \theta) = \frac{e^{z_{O,ik}}}{\sum_{\ell} e^{z_{O,i\ell}}}$$

$$L(\theta) = - \sum_i r_{ik} \ln P(y_i = k | \mathbf{x}_i, \theta) = \sum_i \ln \left(\sum_k e^{z_{O,ik}} \right) - \sum_k r_{ik} z_{O,ik}$$

◦ Called the **categorical cross-entropy**

Outline

- ❑ Motivating Idea: Nonlinear classifiers from linear features
- ❑ Neural Networks
- ❑ Neural Network Loss Function
-  ❑ Stochastic Gradient Descent
 - ❑ Building and Training a Network in Keras
 - Synthetic data
 - MNIST
- ❑ Backpropagation Training

Problems with Standard Gradient Descent

□ Neural network training (like all training): Minimize loss function

$$\hat{\theta} = \arg \min_{\theta} L(\theta), \quad L(\theta) = \frac{1}{N} \sum_{i=1}^N L_i(\theta, \mathbf{x}_i, y_i)$$


- $L_i(\theta, \mathbf{x}_i, y_i)$ = loss on sample i for parameter θ

□ Standard gradient descent:

$$\theta^{k+1} = \theta^k - \alpha \nabla L(\theta^k) = \theta^k - \frac{\alpha}{N} \sum_{i=1}^N \nabla L_i(\theta^k, \mathbf{x}_i, y_i)$$

- Each iteration requires computing N loss functions and gradients
- Will discuss how to compute later
- But, gradient computation is expensive when data size N large

Stochastic Gradient Descent



□ In each step:

- Select random small “mini-batch”
- Evaluate gradient on mini-batch

□ For $t = 1$ to N_{steps}

- Select random mini-batch $I \subset \{1, \dots, N\}$
- Compute gradient approximation:

$$g^t = \frac{1}{|I|} \sum_{i \in I} \nabla L(x_i, y_i, \theta)$$

- Update parameters:

$$\theta^{t+1} = \theta^t - \alpha^t g^t$$

SGD Theory (Advanced)

□ Mini-batch gradient = true gradient in expectation:

$$E(g^t) = \frac{1}{N} \sum_{i=1}^N \nabla L(x_i, y_i, \theta) = \nabla L(\theta^t)$$

□ Hence can write $g^t = \nabla L(\theta^t) + \xi^t$,

- ξ^t = random error in gradient calculation, $E(\xi^t) = 0$
- SGD update: $\theta^{t+1} = \theta^t - \alpha^t g^t$, $\theta^{t+1} = \theta^t - \alpha^t \nabla L(\theta^t) - \alpha^t \xi^t$

□ **Robins-Munro**: Suppose that $\alpha^t \rightarrow 0$ and $\sum_t \alpha^t = \infty$. Let $s_t = \sum_{k=0}^t \alpha^k$

- Then $\theta^t \rightarrow \theta(s_t)$ where $\theta(s)$ is the continuous solution to the differential equation:

$$\frac{d\theta(s)}{ds} = -\nabla L(\theta)$$

□ High-level take away:

- If step size is decreased, random errors in sub-sampling are averaged out

SGD Practical Issues

□ Terminology:

- Suppose **minibatch** size is B . Training size is N
- Each training **epoch** includes updates going through all non-overlapping minibatches
- There are $\frac{N}{B}$ **steps** per training **epoch**

□ Data shuffling

- Generally do not randomly pick a mini-batch
- In each epoch, randomly shuffle training samples
- Then, select mini-batches in order through the shuffled training samples.
- **It is critical to reshuffle in each epoch!**

Outline

- ❑ Motivating Idea: Nonlinear classifiers from linear features
- ❑ Neural Networks
- ❑ Neural Network Loss Function
- ❑ Stochastic Gradient Descent
- ❑ Building and Training a Network in Keras
 - ➔ Synthetic data
 - MNIST
- ❑ Backpropagation Training

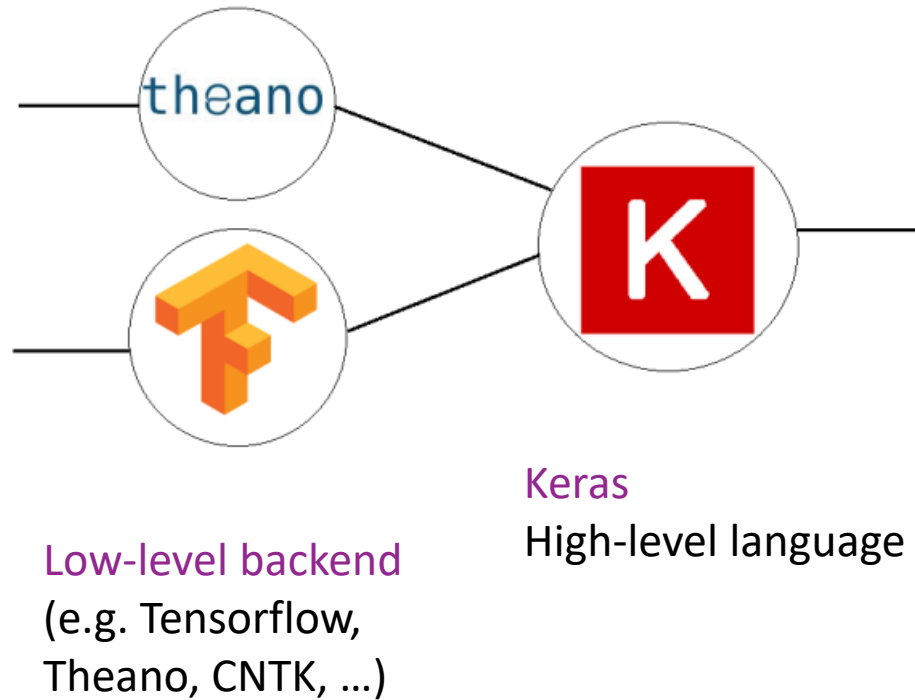
Deep Learning Zoo

- Torch
- Caffe
- Theano (Keras, Lasagne)
- CuDNN
- Tensorflow
- Mxnet
- Etc.



Keras Package

Python Deep Learning Library





- ❑ High-level neural network language in Python
- ❑ Runs on top of a backend
 - Much simpler than raw backend language
 - Very fast coding
 - Uniform language for all backend
- ❑ Keras has been incorporated into TF
- ❑ But...
 - Slightly less flexible
 - Not as fast sometimes
- ❑ In this class, we use Keras

Keras Recipe

- ❑ Step 1. Describe model architecture
 - Number of hidden units, output units, activations, ...
- ❑ Step 2. Select an optimizer
- ❑ Step 3. Select a loss function and compile the model
- ❑ Step 4. Fit the model
- ❑ Step 5. Test / use the model

Synthetic Data Example

- Try a simpler two-layer NN
 - Input $x = 2$ dim
 - 4 hidden units
 - 1 output unit (binary classification)

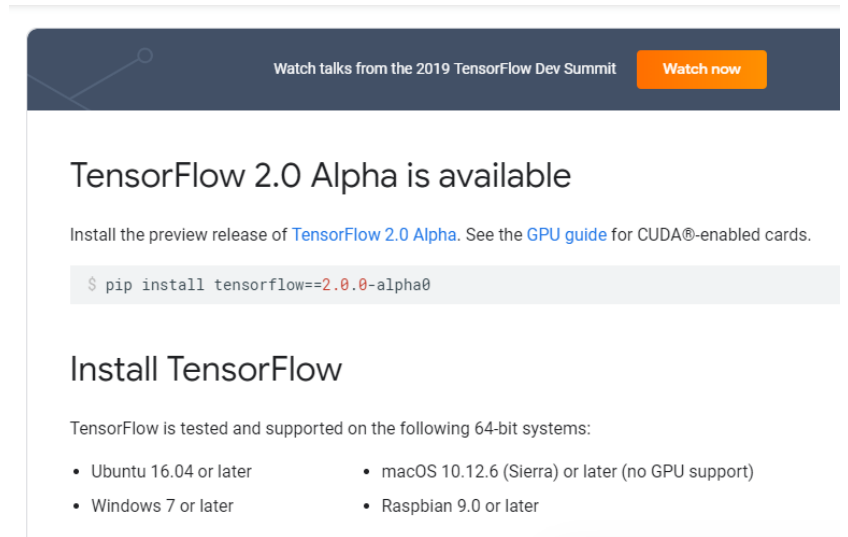


Step 0: Import the Packages

- ❑ Install Tensorflow
- ❑ For this lab, you can use the CPU version
- ❑ If you are using Google Collaboratory, TF is pre-installed

```
import tensorflow as tf
```

<https://www.tensorflow.org/install>



The screenshot shows the TensorFlow website's announcement for TensorFlow 2.0 Alpha. At the top, there is a dark blue header with the text "Watch talks from the 2019 TensorFlow Dev Summit" and an orange "Watch now" button. Below the header, the main content area has a light gray background. The title "TensorFlow 2.0 Alpha is available" is prominently displayed. Underneath, it says "Install the preview release of TensorFlow 2.0 Alpha. See the GPU guide for CUDA®-enabled cards." A code block shows the command: `$ pip install tensorflow==2.0.0-alpha0`. Further down, the section "Install TensorFlow" is followed by a list of supported 64-bit systems: Ubuntu 16.04 or later, macOS 10.12.6 (Sierra) or later (no GPU support), Windows 7 or later, and Raspbian 9.0 or later.

Watch talks from the 2019 TensorFlow Dev Summit [Watch now](#)

TensorFlow 2.0 Alpha is available

Install the preview release of [TensorFlow 2.0 Alpha](#). See the [GPU guide](#) for CUDA®-enabled cards.

```
$ pip install tensorflow==2.0.0-alpha0
```

Install TensorFlow

TensorFlow is tested and supported on the following 64-bit systems:

- Ubuntu 16.04 or later
- macOS 10.12.6 (Sierra) or later (no GPU support)
- Windows 7 or later
- Raspbian 9.0 or later

Step 1: Define Model

```
from tensorflow.keras.models import Model, Sequential  
from tensorflow.keras.layers import Dense, Activation
```

```
import tensorflow.keras.backend as K  
K.clear_session()
```

- ☐ Load modules for layers
- ☐ Clear graph (extremely important!)
- ☐ Build model
 - This example: **dense** layers
 - Give each layer a dimension, name & activation

```
nin = nx # dimension of input data  
nh = 4   # number of hidden units  
nout = 1 # number of outputs = 1 since this is binary  
model = Sequential()  
model.add(Dense(units=nh, input_shape=(nx,), activation='sigmoid', name='hidden'))  
model.add(Dense(units=nout, activation='sigmoid', name='output'))
```


Step 1: Continued

- ❑ Print the model summary
- ❑ For each layers
 - Shows dimensions and shape
- ❑ Note shapes:
 - (None, 4)

Batch size

This is not fixed

Dim per sample in batch

```
model.summary()
```

| Layer (type) | Output Shape | Param # |
|----------------|--------------|---------|
| hidden (Dense) | (None, 4) | 12 |
| output (Dense) | (None, 1) | 5 |

Total params: 17
Trainable params: 17
Non-trainable params: 0

Step 2, 3: Select an Optimizer & Compile

```
from tensorflow.keras import optimizers

opt = optimizers.Adam(lr=0.01)
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

- ❑ Adam optimizer generally works well for most problems
 - In this case, had to manually set learning rate
 - You often need to play with this.
- ❑ Use binary cross-entropy loss
- ❑ Metrics indicate what will be printed in each epoch

Step 4: Fit the Model

```
model.fit(X, y, epochs=10, batch_size=100)
```

```
Epoch 1/10  
400/400 [=====] - 0s - loss: 0.8047 - acc: 0.3900  
Epoch 2/10  
400/400 [=====] - 0s - loss: 0.7695 - acc: 0.3900  
Epoch 3/10  
400/400 [=====] - 0s - loss: 0.7428 - acc: 0.3900  
Epoch 4/10  
400/400 [=====] - 0s - loss: 0.7223 - acc: 0.3900  
Epoch 5/10  
400/400 [=====] - 0s - loss: 0.7027 - acc: 0.4000  
Epoch 6/10  
400/400 [=====] - 0s - loss: 0.6895 - acc: 0.5650  
Epoch 7/10  
400/400 [=====] - 0s - loss: 0.6814 - acc: 0.6100  
Epoch 8/10  
400/400 [=====] - 0s - loss: 0.6756 - acc: 0.6100  
Epoch 9/10  
400/400 [=====] - 0s - loss: 0.6720 - acc: 0.6100  
Epoch 10/10  
400/400 [=====] - 0s - loss: 0.6694 - acc: 0.6100
```

❑ Use keras fit function

- Specify number of epoch & batch size

❑ Prints progress after each epoch

- Loss = loss on training data
- Acc = accuracy on training data

Fitting the Model with Many Epochs

- ❑ This example requires large number of epochs (~1000)
- ❑ Do not want to print progress on each epoch
- ❑ Rewrite code to manually print progress
- ❑ Can also use a **callback** function

```
epoch= 50 loss= 6.6854e-01 acc=0.61000
epoch= 100 loss= 6.6702e-01 acc=0.61000
epoch= 150 loss= 6.5264e-01 acc=0.61000
epoch= 200 loss= 5.9691e-01 acc=0.53500
epoch= 250 loss= 5.4305e-01 acc=0.70500
epoch= 300 loss= 4.8620e-01 acc=0.79000
epoch= 350 loss= 4.1364e-01 acc=0.86250
epoch= 400 loss= 3.6114e-01 acc=0.86250
epoch= 450 loss= 3.3093e-01 acc=0.86750
epoch= 500 loss= 3.1383e-01 acc=0.86750
epoch= 550 loss= 3.0321e-01 acc=0.87250
epoch= 600 loss= 2.9631e-01 acc=0.88000
epoch= 650 loss= 2.9159e-01 acc=0.87750
epoch= 700 loss= 2.8804e-01 acc=0.88250
epoch= 750 loss= 2.8534e-01 acc=0.88750
epoch= 800 loss= 2.8322e-01 acc=0.88250
epoch= 850 loss= 2.8132e-01 acc=0.88750
epoch= 900 loss= 2.7995e-01 acc=0.89000
epoch= 950 loss= 2.7846e-01 acc=0.88500
epoch=1000 loss= 2.7721e-01 acc=0.89000
```

```
nit = 20    # number of training iterations
nepoch_per_it = 50 # number of epochs per iterations

# Loss, accuracy and epoch per iteration
loss = np.zeros(nit)
acc = np.zeros(nit)
epoch_it = np.zeros(nit)


# Main iteration loop
for it in range(nit):

    # Continue the fit of the model
    init_epoch = it*nepoch_per_it
    model.fit(X, y, epochs=nepoch_per_it, batch_size=100, verbose=0)

    # Measure the Loss and accuracy on the training data
    lossi, acci = model.evaluate(X,y, verbose=0)
    epochi = (it+1)*nepoch_per_it
    epoch_it[it] = epochi
    loss[it] = lossi
    acc[it] = acci
    print("epoch=%4d loss=%12.4e acc=%7.5f" % (epochi,lossi,acci))
```

Performance vs Epoch

□ Can observe loss function slowly converging



Step 5. Visualizing the Decision Regions


- ❑ Feed in data $x = (x_1, x_2)$ over grid of points in $[0,1] \times [0,1]$
- ❑ Use predict to observe output for each input point
- ❑ Plot outputs $u_O = \text{sigmoid}(z_O)$

```
# Limits to plot the response.
xmin = [0,0]
xmax = [1,1]

# Use meshgrid to create the 2D input
nplot = 100
x0plot = np.linspace(xmin[0],xmax[1],nplot)
x1plot = np.linspace(xmin[0],xmax[1],nplot)
x0mat, x1mat = np.meshgrid(x0plot,x1plot)
Xplot = np.column_stack([x0mat.ravel(), x1mat.ravel()])

# Compute the output
yplot = model.predict(Xplot)
yplot_mat = yplot[:,0].reshape((nplot, nplot))


# Plot the recovered region
plt.imshow(np.flipud(yplot_mat), extent=[xmin[0],xmax[0],xmin[0],xmax[1]], cmap=plt.cm.Reds)
plt.colorbar()
```



Visualizing the Hidden Layers

```
# Get the response in the hidden units
layer_hid = model.get_layer('hidden')
model1 = Model(inputs=model.input,
               outputs=layer_hid.output)
zhid_plot = model1.predict(Xplot)
zhid_plot = zhid_plot.reshape((nplot,nplot,nh))
```


- ❑ Create a new model with hidden layer output
- ❑ Feed in data $x = (x_1, x_2)$ over $[0,1] \times [0,1]$
- ❑ Predict outputs from hidden outputs



Each hidden layer is a logistic regression layer with a different separating line!

In-Class Exercise

□ Go to in class exercise in github



Now try to have a neural network *learn* the relation $y=f(x)$.

- Clear the keras session
- Create a neural network with 4 hidden units, 1 output unit
- Use a sigmoid activation for the hidden units and no output activation
- Compile with `mean_squared_error` for the loss and metrics
- Fit the model
- Plot the predicted and true function

Outline

- ❑ Motivating Idea: Nonlinear classifiers from linear features
- ❑ Neural Networks
- ❑ Neural Network Loss Function
- ❑ Stochastic Gradient Descent
- ❑ Building and Training a Network in Keras
 - Synthetic data
- ➡ MNIST
- ❑ Backpropagation Training

Recap: MNIST data

- ❑ Classic MNIST problem:
 - Detect hand-written digits
 - Each image is $28 \times 28 = 784$ pixels
- ❑ Dataset size:
 - 50,000 training digits
 - 10,000 test
 - 10,000 validation (not used here)
- ❑ Can be loaded with sklearn and many other packages



Simple MNIST Neural Network

□ 784 inputs, 100 hidden units, 10 outputs

```
nin = Xtr.shape[1] # dimension of input data
nh = 100          # number of hidden units
nout = int(np.max(ytr)+1) # number of outputs = 10 since there are 10 classes
model = Sequential()
model.add(Dense(units=nh, input_shape=(nin,), activation='sigmoid', name='hidden'))
model.add(Dense(units=nout, activation='softmax', name='output'))
```

```
model.summary()
```

| Layer (type) | Output Shape | Param # |
|----------------|--------------|---------|
| hidden (Dense) | (None, 100) | 78500 |
| output (Dense) | (None, 10) | 1010 |

=====
Total params: 79,510
Trainable params: 79,510
Non-trainable params: 0

Why 78500 parameters in hidden layer?

Fitting the Model


- ❑ Run for 20 epochs, ADAM optimizer, batch size = 100
- ❑ Final accuracy = 0.972
- ❑ Not great, but much faster than SVM. Also CNNs we study later do even better.

```
opt = optimizers.Adam(lr=0.001) # beta_1=0.9, beta_2=0.999
model.compile(optimizer=opt,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
model.fit(Xtr, ytr, epochs=10, batch_size=100, validation_data=(Xts,yts))
```

```
Epoch 1/10
50000/50000 [=====] - 3s - loss: 0.0474 - acc: 0.9868 - val_loss: 0.0886 - val_ac
c: 0.9717
Epoch 8/10
50000/50000 [=====] - 3s - loss: 0.0440 - acc: 0.9884 - val_loss: 0.0875 - val_ac
c: 0.9718
Epoch 9/10
50000/50000 [=====] - 2s - loss: 0.0393 - acc: 0.9903 - val_loss: 0.0872 - val_ac
c: 0.9732
Epoch 10/10
50000/50000 [=====] - 3s - loss: 0.0381 - acc: 0.9901 - val_loss: 0.0875 - val_ac
c: 0.9718
```

Training and Validation Accuracy




```
tr_accuracy = hist.history['acc']
val_accuracy = hist.history['val_acc']

plt.plot(tr_accuracy)
plt.plot(val_accuracy)
plt.grid()
plt.xlabel('epochs')
plt.ylabel('accuracy')
plt.legend(['training accuracy', 'validation accuracy'])
```

- Training accuracy continues to increase
- Validation accuracy eventually flattens and sometimes starts to decrease.
- Should stop when the validation accuracy starts to decrease.
- This indicates overfitting.

Outline

- ❑ Motivating Idea: Nonlinear classifiers from linear features
- ❑ Neural Networks
- ❑ Neural Network Loss Function
- ❑ Stochastic Gradient Descent
- ❑ Building and Training a Network in Keras
 - Synthetic data
 - MNIST
-  ❑ Backpropagation Training

Stochastic Gradient Descent

□ Training uses SGD

□ In each step:

- Select a subset of sample for minibatch $I \subset \{1, \dots, N\}$
- Evaluate mini-batch loss $L(\theta^t) = \sum_{i \in I} L_i(\theta^t, \mathbf{x}_i, y_i)$
- Evaluate mini-batch gradient $\mathbf{g}^t = \sum_{i \in I} \nabla L_i(\theta^t, \mathbf{x}_i, y_i)$
- Take SGD step: $\theta^{t+1} = \theta^t - \alpha \mathbf{g}^t$

□ Question: How do we compute gradient?

Gradients with Multiple Parameters

- For neural net problem: $\theta = (W_H, b_H, W_o, b_o)$
- Gradient is computed with respect to each parameter:


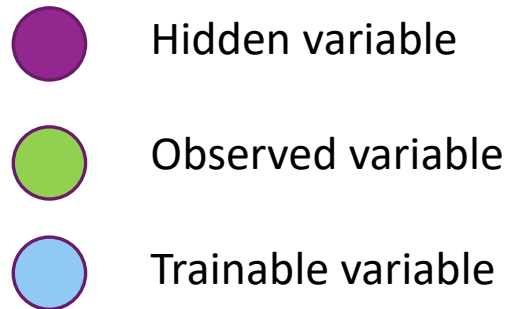
$$\nabla L(\theta) = [\nabla_{W_H} L(\theta), \nabla_{b_H} L(\theta), \nabla_{W_o} L(\theta), \nabla_{b_o} L(\theta)]$$

- Gradient descent is performed on each parameter:

$$\begin{aligned} W_H &\leftarrow W_H - \alpha \nabla_{W_H} L(\theta), \\ b_H &\leftarrow b_H - \alpha \nabla_{b_H} L(\theta), \\ &\dots \end{aligned}$$

Computation Graph & Forward Pass

- ❑ Neural network loss function can be computed via a **computation graph**
- ❑ Sequence of operations starting from measured data and parameters
- ❑ Loss function computed via a **forward pass** in the computation graph
 - $z_{H,i} = W_H x_i + b_H$
 - $u_{H,i} = g_{act}(z_{H,i})$
 - $z_{O,i} = W_O u_{H,i} + b_O$
 - $L = \sum_i L_i(z_{O,i}, y_i)$



Forward Pass Example in Numpy

□ Example network:

- Single hidden layer with N_H hidden units, single output unit
- Sigmoid activation, binary cross entropy loss

```
def loss(X,y,theta):  
    """  
    Computes loss function for neural network  
    with sigmoid activation, binary cross-entropy loss  
    """  
    # Unpack parameters  
    Wh, bh, Wo, bo = theta  
  
    # Hidden layer  
    Zh = X.dot(Wh) + bh[None,:]  
    Uh = 1/(1+np.exp(-Zh))  
  
    # Output layer  
    Zout = Uh.dot(Wo) + bo[None,:]  
    Uout = 1/(1+np.exp(-Zout))  
  
    # Loss function  
    f = np.sum(-y*Zout + np.log(1+y*Zout))  
    return f
```

```
nh = 4      # number hidden units  
nin = 2     # input dimension  
nout = 1    # output dimension  
nsamp = 100 # number samples in batch  
  
# Random data  
X = np.random.randn(nsamp,nin)/np.sqrt(nin)  
y = (np.random.rand(nsamp) < 0.5).astype(float)  
  
# Random weights  
Wh = np.random.randn(nin,nh)  
bh = np.random.randn(nh)  
Wo = np.random.randn(nh,nout)  
bo = np.random.randn(nout)  
  
# Compute Loss  
f = loss(X,y,[Wh,bh,Wo,bo])
```

Back-Propagation on A Two Node Graph

□ Back Propagation:

- A way to compute gradients
- Iterative procedure that works in reverse

□ Consider a simple 2 node computation graph

□ Each component is a vector:


- $\mathbf{x} = (x_1, \dots, x_N)$, $\mathbf{z} = (z_1, \dots, z_M)$

□ First we compute $\frac{\partial J}{\partial z_i}$

□ Then compute $\frac{\partial J}{\partial x_j}$ from multi-variable chain rule:

$$\frac{\partial J}{\partial x_j} = \sum_{i=1}^n \frac{\partial J}{\partial z_i} \frac{\partial z_i}{\partial x_j}$$

Variables computed in forward pass



Gradients computed in reverse pass

Back-Prop on a General Computation Graph

□ Backpropagation:


- Compute gradients backwards
- Work one node at a time

□ First compute all derivatives of all the variables

- $\partial L / \partial z_O$
- $\partial L / \partial u_H$ from $\partial L / \partial z_O, \partial z_O / \partial u_H$
- $\partial L / \partial z_H$ from $\partial L / \partial u_H, \partial u_H / \partial z_H$

□ Then compute gradient of parameters:

- $\partial L / \partial W_O$ from $\partial L / \partial z_O, \partial z_O / \partial W_O$
- $\partial L / \partial b_O$ from $\partial L / \partial z_O, \partial z_O / \partial b_O$
- $\partial L / \partial W_H$ from $\partial L / \partial z_H, \partial z_H / \partial W_H$
- $\partial L / \partial b_H$ from $\partial L / \partial z_H, \partial z_H / \partial b_H$
-




Back-Propagation Example (Part 1)

Continue our example:

- Single hidden layer with M hidden units, single output unit
- Sigmoid activation, binary cross entropy loss
- N samples, D input dimension

Loss node:

- $L = \sum_i L_i(z_{O,i}, y_i)$
- $L_i(z_{oi}) = -(y_i \ln P(y_i = 1|x_i, \theta) + (1 - y_i) \ln P(y_i = 0|x_i, \theta)) = y_i z_{oi} - \ln[1 + e^{y_i z_{oi}}]$
- $\frac{\partial L}{\partial z_{O,i}} = y_i - \frac{e^{y_i z_{oi}}}{1 + e^{y_i z_{oi}}}$



Back-Propagation Example (Part 2)

□ Node z_O


- $z_O = u_H W_O + b_O$
- $z_{O,i} = \sum_m u_{H,im} W_{Om} + b_O$

□ Gradient:

- $\frac{\partial z_{O,i}}{\partial W_{O,m}} = u_{H,i,m}$
- $\frac{\partial z_{O,i}}{\partial b_O} = 1$
- Other partial derivatives are zero

□ Apply chain rule:

- $\frac{\partial L}{\partial W_{O,m}} = \sum_i \frac{\partial L}{\partial z_{O,i}} \frac{\partial z_{O,i}}{\partial W_{O,m}} = \sum_i \frac{\partial L}{\partial z_{O,i}} u_{H,i,m}$
- $\frac{\partial L}{\partial b_O} = \sum_i \frac{\partial L}{\partial z_{O,i}} \frac{\partial z_{O,i}}{\partial b_O} = \sum_i \frac{\partial L}{\partial z_{O,i}}$



Back-Propagation Example (Part 3)

□ Node z_O


- $z_O = u_H W_O + b_O$
- $z_{O,i} = \sum_m u_{H,im} W_{Om} + b_O$

□ Gradient:

- $\frac{\partial z_{O,i}}{\partial u_{H,ij}} = W_{O,j}$, $m=1, \dots, M$
- Other partial derivatives are zero


□ Apply chain rule:

- $\frac{\partial L}{\partial u_{H,ij}} = \frac{\partial L}{\partial z_{O,i}} \frac{\partial z_{O,i}}{\partial u_{H,ij}} = \frac{\partial L}{\partial z_{O,i}} W_{O,j}$



Back-Propagation Example (Part 4,...)

- Compute $\frac{\partial L}{\partial W_H}, \frac{\partial L}{\partial b_H}$
- Following previous procedures, in class



- Summary:
 - Forward pass: Compute hidden nodes and loss
 - Backward pass: Compute gradients

Initialization and Data Normalization

- ❑ Solution by gradient descent algorithm depends on the initial solution
- ❑ Typically weights are set to random values near zero.
- ❑ Small weights make the network behave like linear classifier.
 - Hence model starts out nearly linearly
 - Becomes nonlinear as weights increase during the training process.
- ❑ Starting with large weights often lead to poor results.
- ❑ Normalizing data to zero mean and unit variance
 - Allows all input dimensions be treated equally and facilitate better convergence.
- ❑ With normalized data, it is typical to initialize the weights to be uniform in $[-0.7, 0.7]$ [ESL]

Regularization

❑ To avoid the weights get too large, can add a penalty term explicitly, with regularization level λ

❑ Ridge penalty

$$R(\theta) = \sum_{d,m} w_{H,d,m}^2 + \sum_{m,k} w_{O,m,k}^2 = \|w_H\|^2 + \|w_O\|^2$$

❑ Total loss

$$L_{reg}(\theta) = L(\theta) + \lambda R(\theta)$$

❑ Change in gradient calculation

❑ Typically used regularization

- L2 = Ridge: Shrink the sizes of weights
- L1: Prefer sparse set of weights
- L1-L2: use a combination of both

Regularization in Keras

- `kernel_regularizer`: instance of `keras.regularizers.Regularizer`
- `bias_regularizer`: instance of `keras.regularizers.Regularizer`
- `activity_regularizer`: instance of `keras.regularizers.Regularizer`

Activity regularization tries to make the output at each layer small or sparse.

Example

```
from keras import regularizers
model.add(Dense(64, input_dim=64,
                  kernel_regularizer=regularizers.l2(0.01),
                  activity_regularizer=regularizers.l1(0.01)))
```

Available penalties

```
keras.regularizers.l1(0.)
keras.regularizers.l2(0.)
keras.regularizers.l1_l2(0.)
```

Choice of network parameters

- ☐ Number of layers (typically not more than 2)
- ☐ Number of hidden units in the hidden layer
- ☐ Regularization level
- ☐ Learning rate
- ☐ Determined by maximizing the cross validation error through typically exhaustive search

Learning Objectives

- ☐ Mathematically describe a neural network with a single hidden layer
 - Describe mappings for the hidden and output units
- ☐ Manually compute output regions for very simple networks
- ☐ Select the loss function based on the problem type
- ☐ Build and train a simple neural network in Keras
- ☐ Write the formulas for gradients using backpropagation
- ☐ Describe mini-batches in stochastic gradient descent
- ☐ Importance of regularization
- ☐ Hyperparameter optimization

Lab for this unit

- ☐ Music instrument classification based on music signals
- ☐ Use hand-crafted features for audio (MFCC)
- ☐ Train a neural net
- ☐ Optimize the learning rate