

Introduction to Machine Learning

Problems: Principal Component Analysis

Profs. Sundeep Rangan and Yao Wang

1. Assume that you have 4 samples each with dimension 3, described in the data matrix X ,

$$X = \begin{bmatrix} 3 & 2 & 1 \\ 2 & 4 & 5 \\ 1 & 2 & 3 \\ 0 & 2 & 5 \end{bmatrix}$$

For the problems below, you may do the calculations in python. Describe the commands you used.

- (a) Find the sample mean.
 - (b) Find the sample covariance matrix Q .
 - (c) Find the eigenvalues and eigenvectors. You can use the `numpy.linalg.eig` to compute eigenvalues and eigenvectors from Q .
 - (d) Find the PCA coefficients corresponding to samples in X .
 - (e) Reconstruct the original samples from the PCA coefficients.
 - (f) Approximate the samples using principle components corresponding to the two largest eigenvalues.
 - (g) Verify the sum of reconstruction error squares = sum of squares of skipped PCA coefficients.
2. *PC bases.* Suppose that the mean and first two PCs in a basis are,

$$\boldsymbol{\mu} = [1, 0, 2], \quad \mathbf{v}_1 = \frac{1}{\sqrt{2}}[1, 1, 0], \quad \mathbf{v}_2 = \frac{1}{\sqrt{2}}[1, -1, 0],$$

- (a) What are the two PC coefficients of the vector $\mathbf{x} = [2, 3, 4]$?
 - (b) What is $\hat{\mathbf{x}}$, the approximation of \mathbf{x} using two PCs?
 - (c) What is the approximation error $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$?
3. *Fitting data with PCs.* You are given python functions for PCA transform and a binary classifier:

```
mu, V = PCA(X) # Finds the mean and PCs for a data matrix X

clf = Classifier()
clf.fit(Z,y)    # Fits a binary classifier from features Z and labels y
yhat = clf.predict(Z) # Predicts labels from Z
```

Given a data matrix \mathbf{x} with binary labels \mathbf{y} , you wish to build a classifier that uses a PCA transform followed by the `Classifier`. Write python code that uses model validation to select the optimal number of PCA components to use. You may assume:

- You use simple cross validation with one training and test split. Not You do not need to do K -fold validation.
 - You should use roughly 25% of the samples for test. Data shuffling before splitting is not needed.
4. *PC with images.* A dataset has 1000 images of size 28×28 , represented as python array \mathbf{x} with shape `(1000,28,28)`. You are given the following python functions:

```
Y = reshape(X, shape)          # Reshapes X to a shape
pca = PCA(n_components = nc)    # Creates a PCA object
pca.fit(Y)                      # Finds the mean and PCs from training data Y
Z = pca.transform(Y)           # Find coefficients in the PC basis
Yhat = pca.inverse_transform(Z) # Invert the PC transform
```

Write a few lines of python code to (i) fit a PC model from the first 500 images with 5 components; and (ii) create an array of approximations of the remaining 500 images.

5. *PCs using SVDs.* You are given a python function:

```
U,s, Vtr = svd(Z, full_matrices = False)
```

which computes an “economy” SVD. Given a data matrix \mathbf{x} , write python code that:

- Finds the PCs and mean of the data.
- Finds the minimum number of PCs in order that the proportion of variance is at least 90%.
- Create an approximation $\hat{\mathbf{x}}$ of \mathbf{x} using the number of PCs in part (ii).